

AI-Powered Nutrition Analyzer For Fitness Enthusiasts

Date	
Team ID	SWTID1749893823
Project Title	AI-Powered Nutrition Analyzer For Fitness Enthusiasts
Maximum Marks	6 Marks

Model Optimization and Tuning Phase

The Model Optimization and Tuning Phase involves refining neural network models for peak performance. It includes optimized model code, fine-tuning hyperparameters, comparing performance metrics, and justifying the final model selection for enhanced predictive accuracy and efficiency.

Hyperparameter Tuning Documentation (8 Marks):

Model	Tuned Hyperparameters
VGG16	<p>In the model, include_top=False removes VGG16’s original classifier to allow custom layers. base_model.trainable=False freezes pre-trained layers to retain learned features. Dropout(0.3) prevents overfitting, and Dense(128) controls model complexity.</p> <pre>from tensorflow.keras.applications import VGG16  # Load VGG16 without top layer base_model = VGG16(weights='imagenet', include_top=False, input_shape=(224, 224, 3)) base_model.trainable = False # Freeze base  # Add classification head model = Sequential([     base_model,     GlobalAveragePooling2D(),     Dropout(0.3),     Dense(128, activation='relu'),     Dropout(0.3),     Dense(5, activation='softmax') ])  model.compile(optimizer=Adam(), loss='categorical_crossentropy', metrics=['accuracy']) model.summary()</pre>
MobileNetV2	<p>Here’s a short paragraph explaining the key parameters in your MobileNetV2-based model:</p> <p>In this model, include_top=False removes MobileNetV2’s original classification layers so custom ones can be added. Setting base_model.trainable=False freezes the pre-trained layers to keep learned features intact. Dropout(0.3) helps reduce overfitting, while Dense(128) defines the custom layer’s complexity for learning fruit classes effectively.</p>

	<pre> import tensorflow as tf from tensorflow.keras.applications import MobileNetV2 from tensorflow.keras.models import Sequential from tensorflow.keras.layers import Dense, Dropout, GlobalAveragePooling2D from tensorflow.keras.optimizers import Adam  # Number of classes (fruits) NUM_CLASSES = 5  # Load MobileNetV2 without the top layer (include_top=False) base_model = MobileNetV2(weights='imagenet', include_top=False, input_shape=(224, 224, 3)) base_model.trainable = False # Freeze the base model  # Build the model model = Sequential([     base_model,     GlobalAveragePooling2D(),     Dropout(0.3),     Dense(128, activation='relu'),     Dropout(0.3),     Dense(NUM_CLASSES, activation='softmax'), ])  # Compile the model model.compile(optimizer=Adam(),               loss='categorical_crossentropy',               metrics=['accuracy'])  model.summary() </pre>
--	--

### Final Model Selection Justification (2 Marks):

Final Model	Reasoning
MobileNetV2	This model was chosen as the final optimized model due to its lightweight architecture, faster training time, and strong performance on image classification tasks with limited computational resources.